

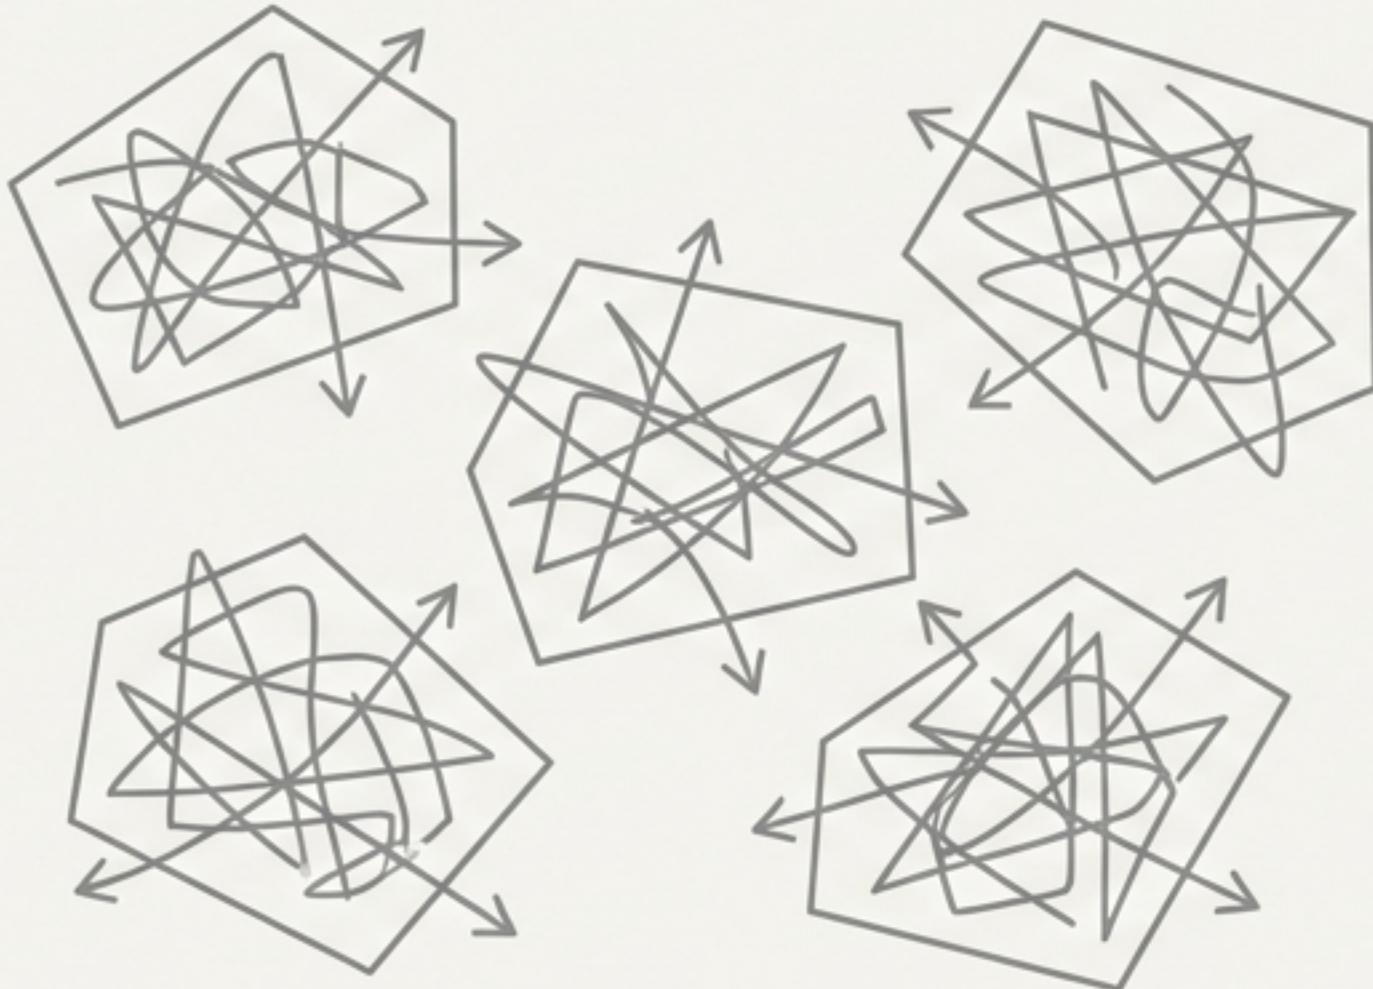


ROS 2: The Operating System for Robots

A Foundational Guide to the Standard Framework for Modern Robotics

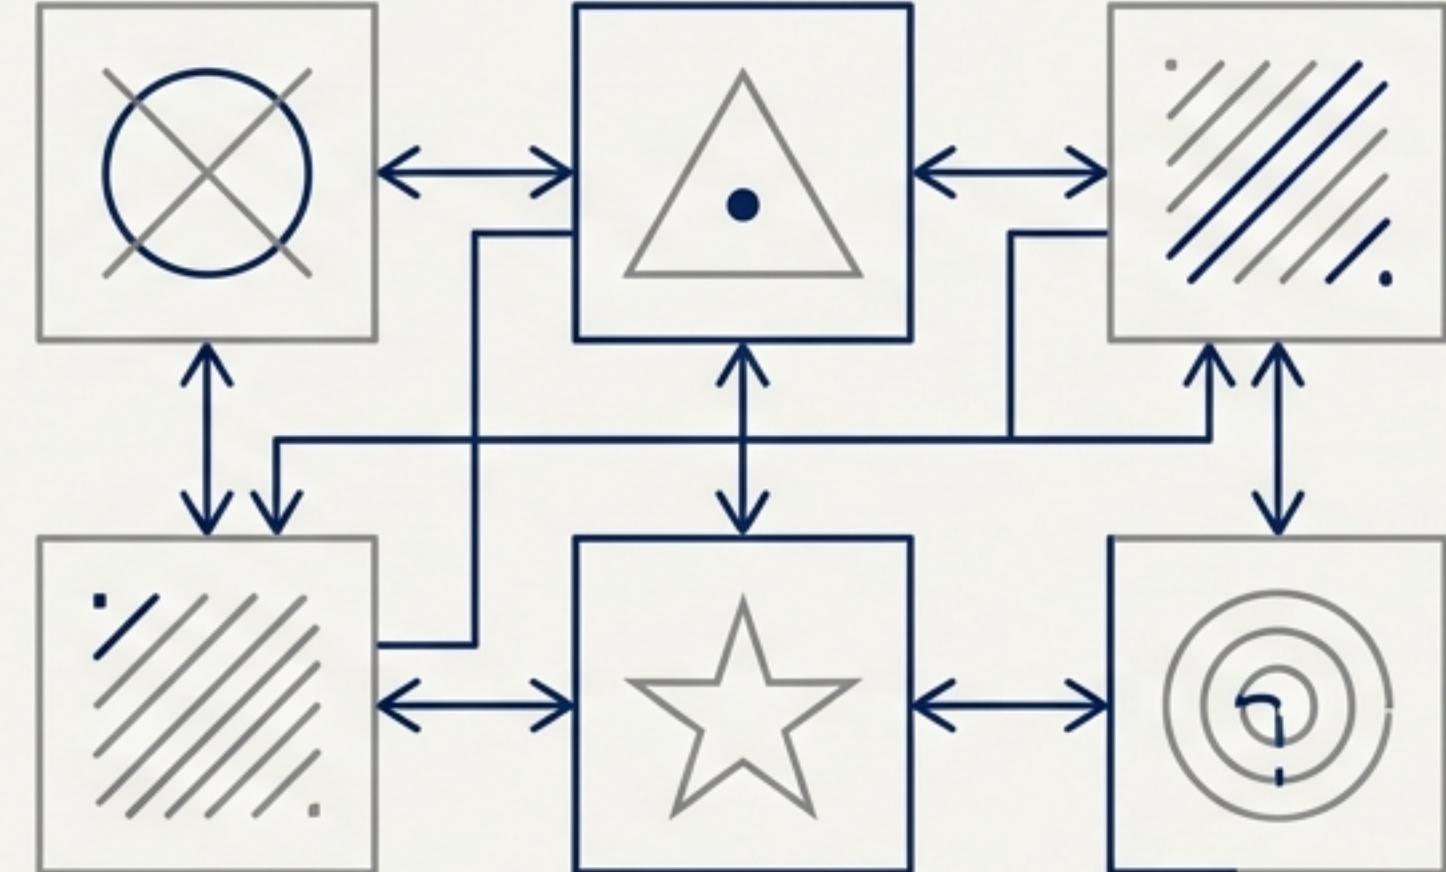
The Core Challenge in Robotics: Stop Reinventing the Wheel

The Siloed Approach



A lot of robotics research effort was spent reinventing the proverbial wheel, especially in software. Custom, non-reusable code for every project leads to fragmentation, integration challenges, and slower progress.

The ROS Framework



An open-source software framework that could be used in a variety of situations and on a variety of hardware. In essence, a Linux for robotics. ROS provides a standardized structure for communication, tools, and capabilities.

ROS is Middleware for Your Robot

The Robot Operating System (ROS) is not a true operating system... rather it is a middleware that sits between your underlying operating system and the robot application code. It's a collection of software tools, libraries, and messaging systems that help break robot code into manageable chunks.

Your Application Code (Perception, Control, Navigation)



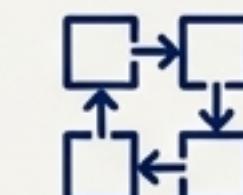
Open-Source Framework

ROS 2 Middleware (Messaging, Tools, Libraries)



Rich Tooling & Debugging

Kernel / OS (Typically Ubuntu Linux)



Manageable, Reusable Code

Hardware (Sensors, Motors, etc.)

From Stanford Lab to Industry Standard



Stanford University.
PhD students propose
an open-source
framework to avoid
“reinventing the wheel.”



Willow Garage.
Founder Scott Hassan
incubates the project.
PR2 prototype proves
basic navigation.



ROS 1.0 Released.
First major public
version.



OSRF Founded.
The Open-Source
Robotics Foundation
becomes the official
maintainer.



ROS 2 Released.
A complete rewrite to
support the needs of
industrial and
commercial robotics.



2007



2008



2010



2012



2017

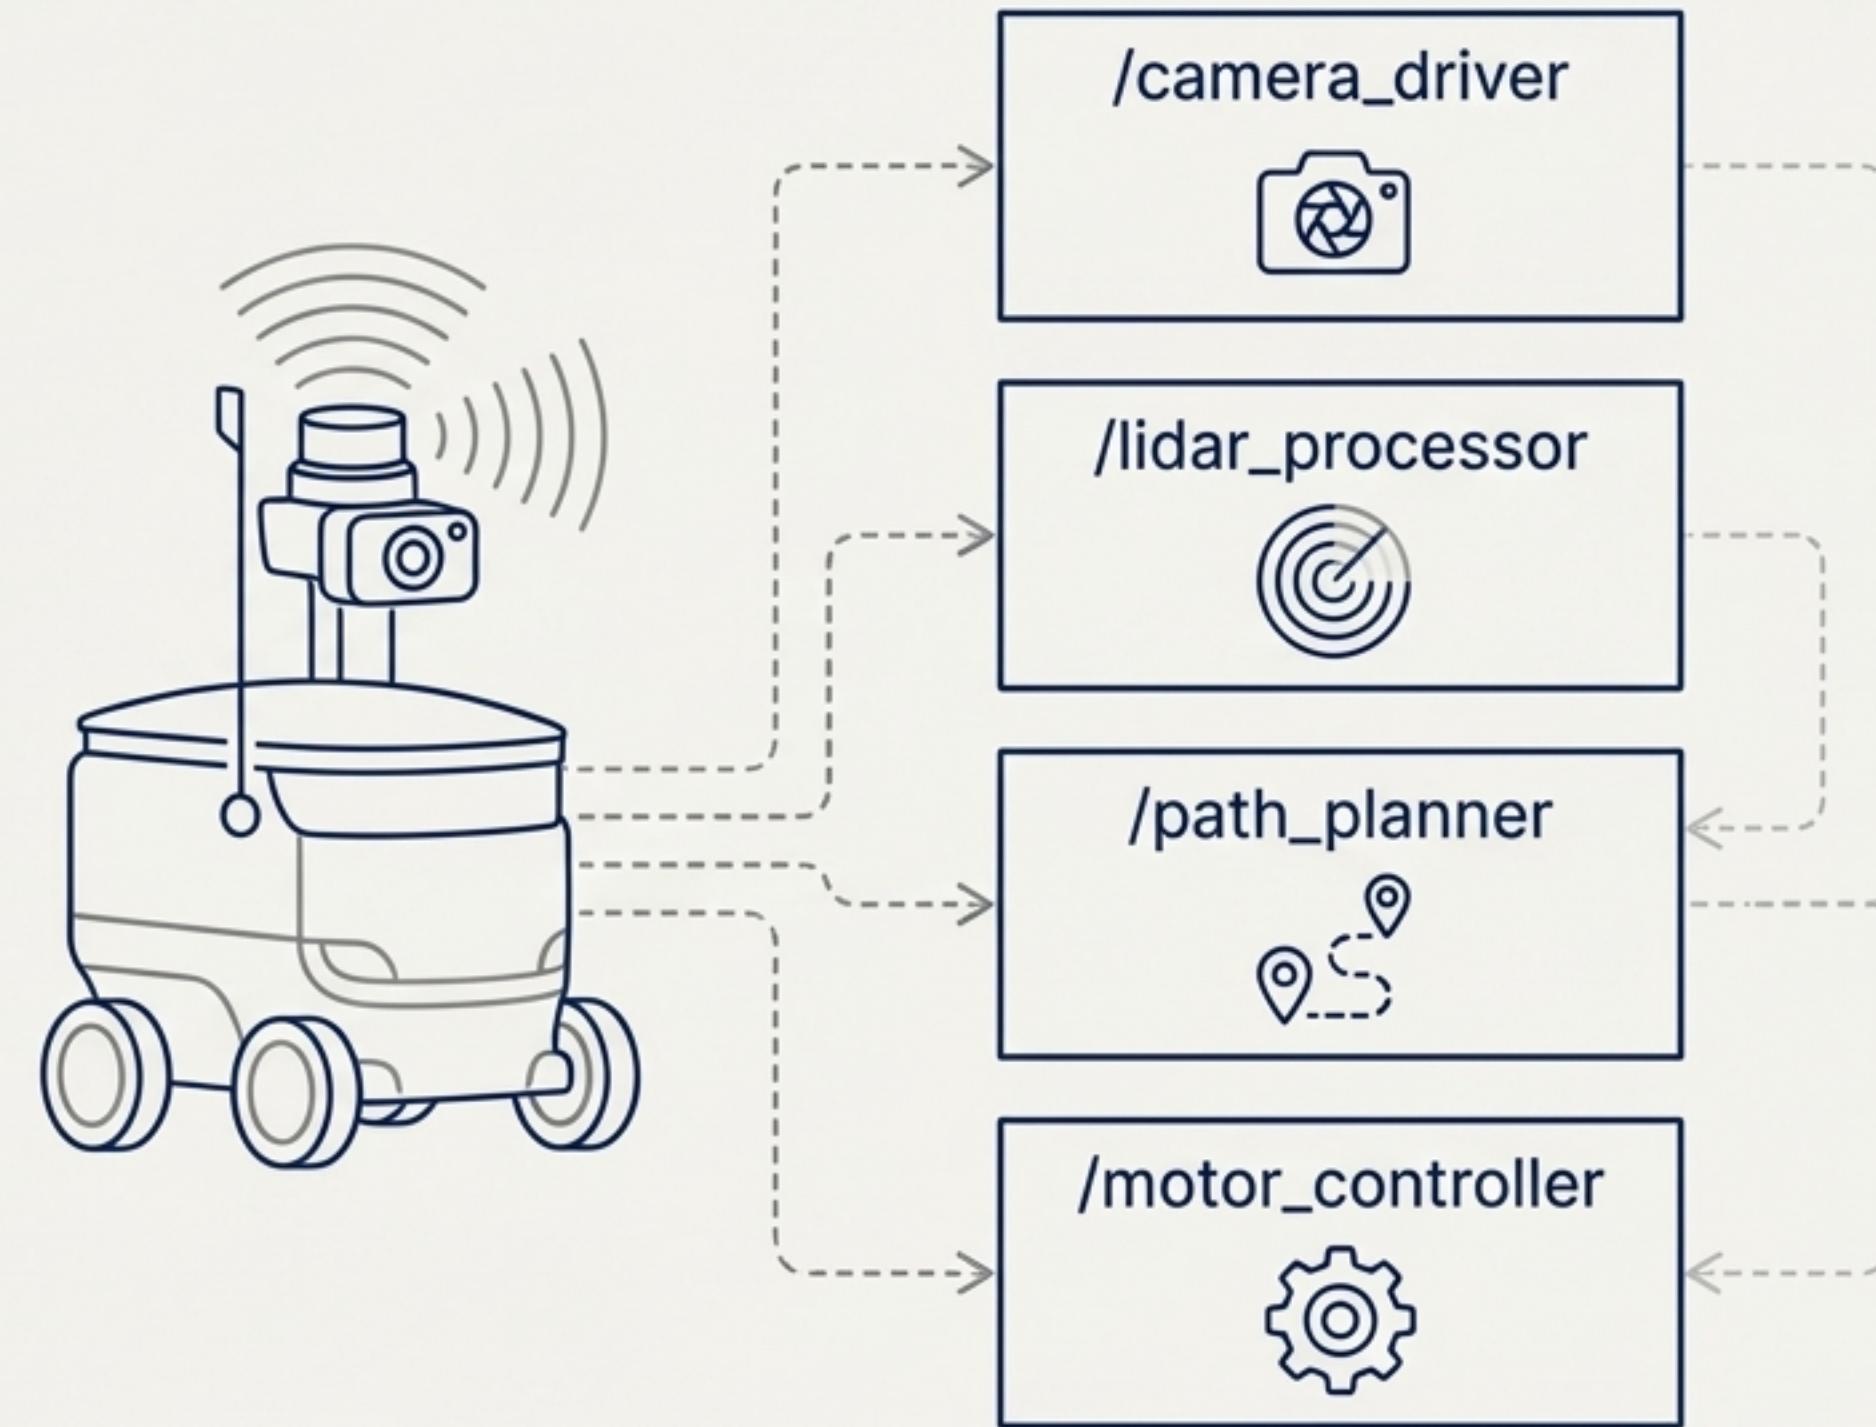
Why ROS 2?

The collective developer team recognized the architectural limitations in the original ROS. ROS 1 will have reached end of life... as a result, we will stick with ROS 2.

The Core Building Block: The ROS Node

In ROS, the fundamental software building block is called a node. This is a process that performs some task, like reading data from a sensor or controlling motors.

- Nodes are typically standalone executables.
- Each node runs in its own process, improving robustness.
- This structure encourages splitting complex systems into separate, reusable files for readability and debugging.

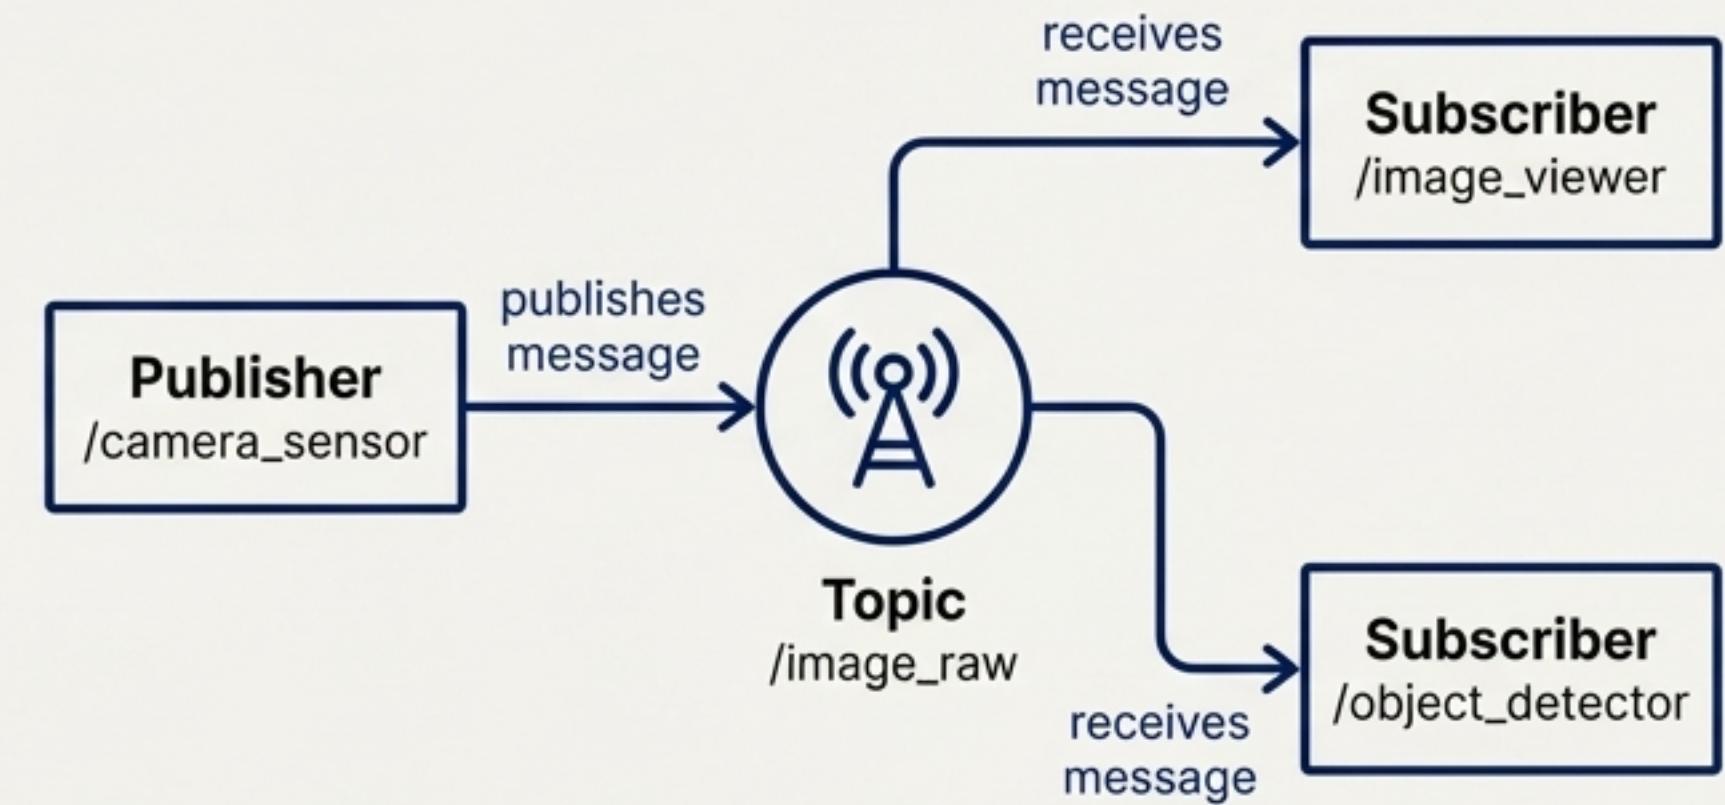


Communication Pattern 1: Topics (Publish/Subscribe)

Topics are named communication channels used for the publish/subscribe model. This is an asynchronous, one-to-many communication method.

How it Works

1. A node, the **Publisher**, sends data to a topic.
2. Any number of nodes, the **Subscribers**, can listen to that topic.
3. Whenever a message is published, all subscribing nodes receive it.

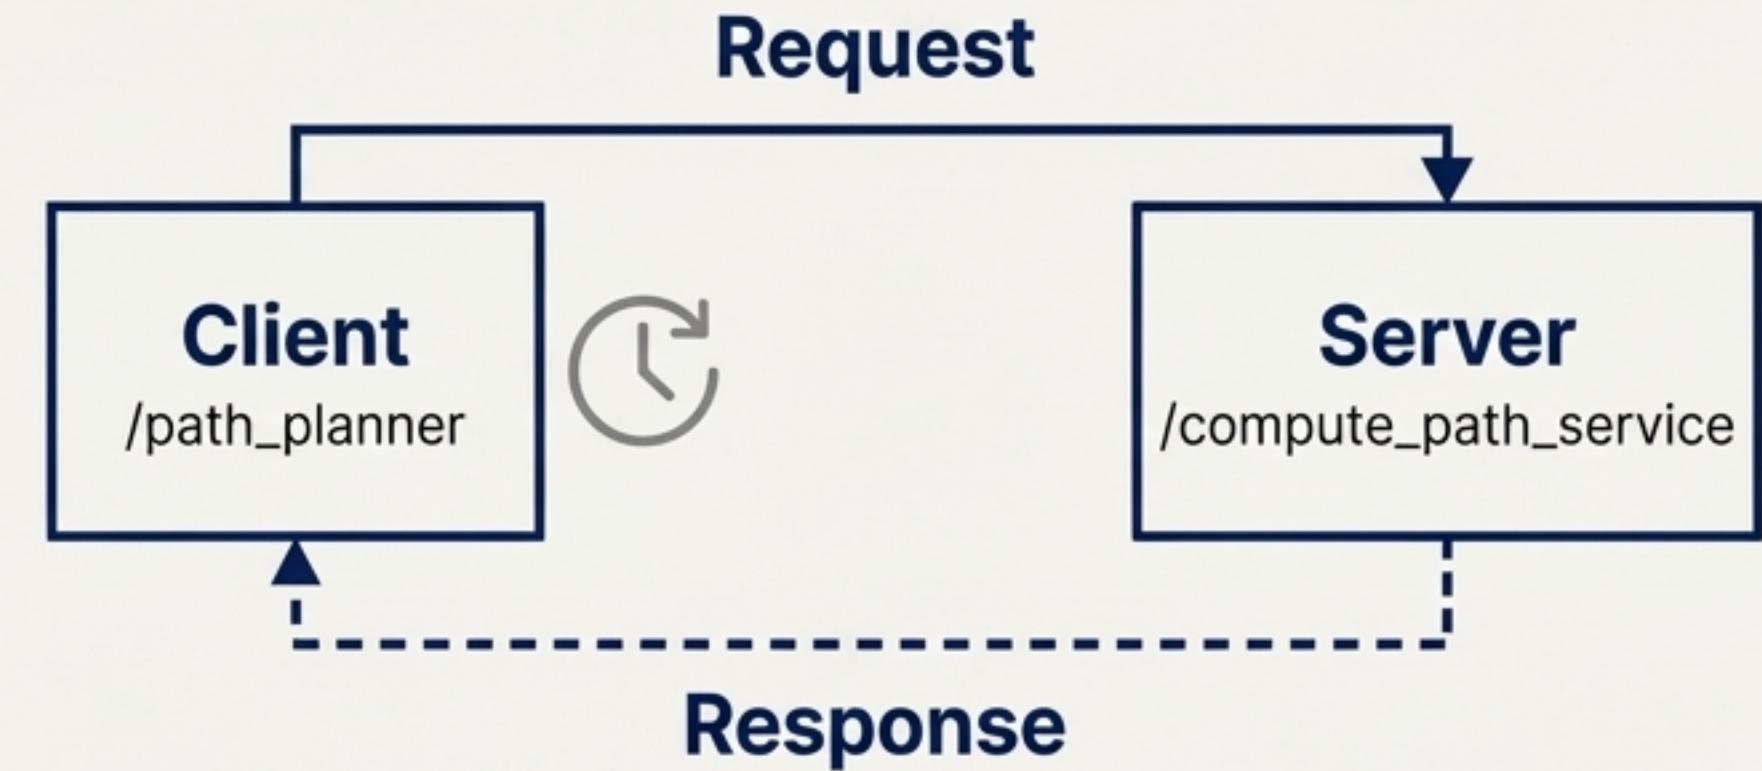


Communication Pattern 2: Services (Request/Response)

Description: A service is a synchronous communication method. This is a one-to-one pattern where a client sends a request and waits for the server to reply with a response.

How it Works

1. A Client node sends a specific request message to a service.
2. The Client pauses execution and waits.
3. A Server node receives the request, performs a task, and sends back a response message.
4. The Client receives the response and resumes its work.



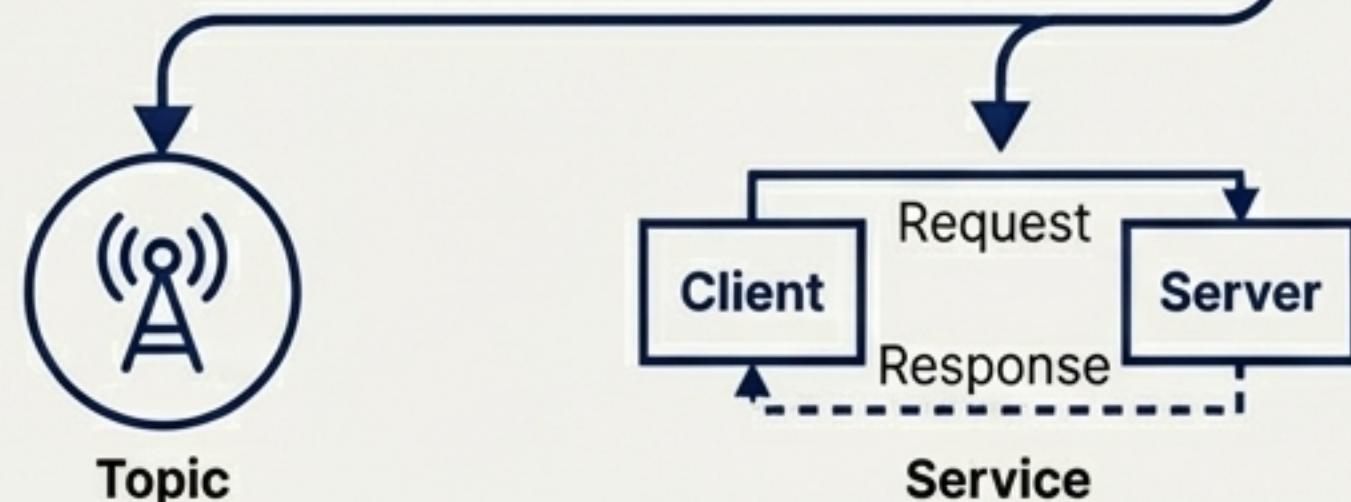
The Universal Language: Interfaces & Messages

Both topics and services rely on strict typing given by interface definitions. Think of it like a programming API: it defines what kind of data is expected to be received or sent by a node.

- **Interfaces** define the format and data structure of messages.
- **Messages** contain one or more fields with strongly typed data (strings, integers, floating-point numbers, arrays, etc.).
- This strict contract allows nodes written in different languages (e.g., Python and C++) to communicate seamlessly.

MyMessage.msg

```
# This defines the structure for a message
string object_name
float64 confidence_score
geometry_msgs/Point position
```



Choosing Your Language: Python vs. C++

They are different tools for different jobs. You can mix and match languages in the same system.



Python (rclpy)

- + **Pro:** Faster Prototyping & Development (Easier to learn and write).
- + **Pro:** Rich Machine Learning Ecosystem (Vast array of AI/ML packages).
- **Con:** Slower Performance (Interpreted language, 10-100x slower than C++).
- **Con:** Higher Memory Overhead & GIL (Global Interpreter Lock limits true multi-threading).
- ⚙️ **Best for:** High-level logic, UI, rapid prototyping, ML integration.



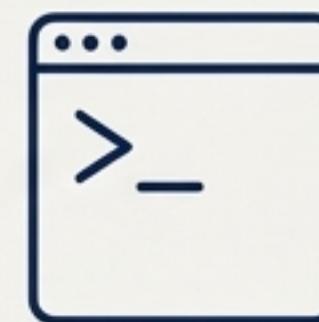
C++ (rclcpp)

- + **Pro:** High Performance (Compiled language for real-time applications).
- + **Pro:** Low-Level Control (Direct memory manipulation).
- + **Pro:** Compile-Time Type Checking (Helps prevent runtime errors).
- **Con:** More Complex Syntax (Verbose, can slow down development).
- ⚙️ **Best for:** Real-time control loops, drivers, performance-critical algorithms.

The ROS 2 Ecosystem is More Than Just Code

Command-Line Tools

A suite of tools (`ros2 run`, `ros2 topic`, etc.) for introspecting, debugging, and running your system.



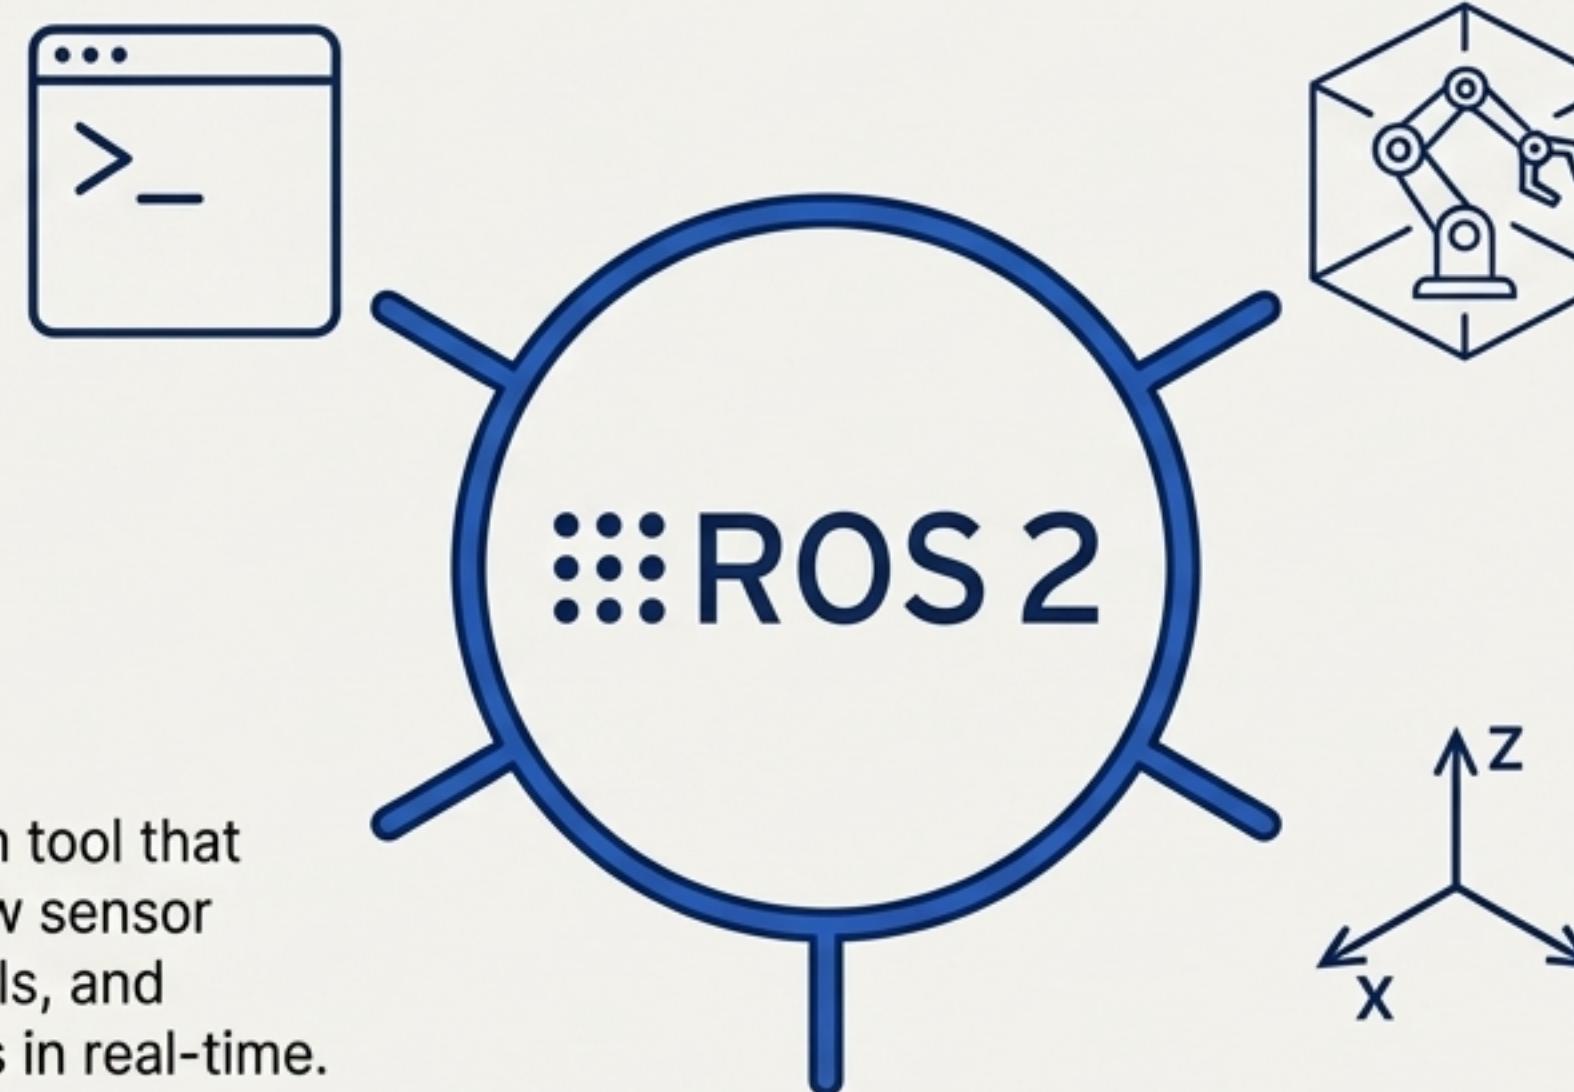
RViz

A 3D visualization tool that allows you to view sensor data, robot models, and algorithm outputs in real-time.



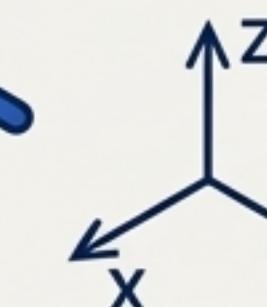
Launch System

A system for starting and configuring multiple nodes with a single command.



Gazebo

A powerful 3D robotics simulator for testing algorithms in realistic environments before deploying to hardware.



TF2 Library

Manages coordinate frames, allowing you to track the position and orientation of multiple moving parts of a robot.

Acknowledging the Limitations

While incredibly powerful, ROS has a few trade-offs to consider. For many applications, its benefits as the de facto industry standard far outweigh them.

Key Considerations

- **Learning Curve:** The abstraction and number of tools can be overwhelming for newcomers.
- **Performance Overhead:** The DDS system and node-based environment can introduce latency.
- **Dependency Management:** Tightly coupled to specific OS versions (usually Ubuntu LTS), which can lead to wrestling with dependencies.

When ROS might not be the best fit

- Small, low-complexity projects.
- Low-power microcontrollers with tight real-time deadlines (though Micro-ROS is an emerging option for this).

Brief Mention of Alternatives

Other frameworks exist (e.g., O-ROCOS, LCM), but ROS is the dominant standard for R&D and prototyping.

Your Recommended Development Environment: Docker

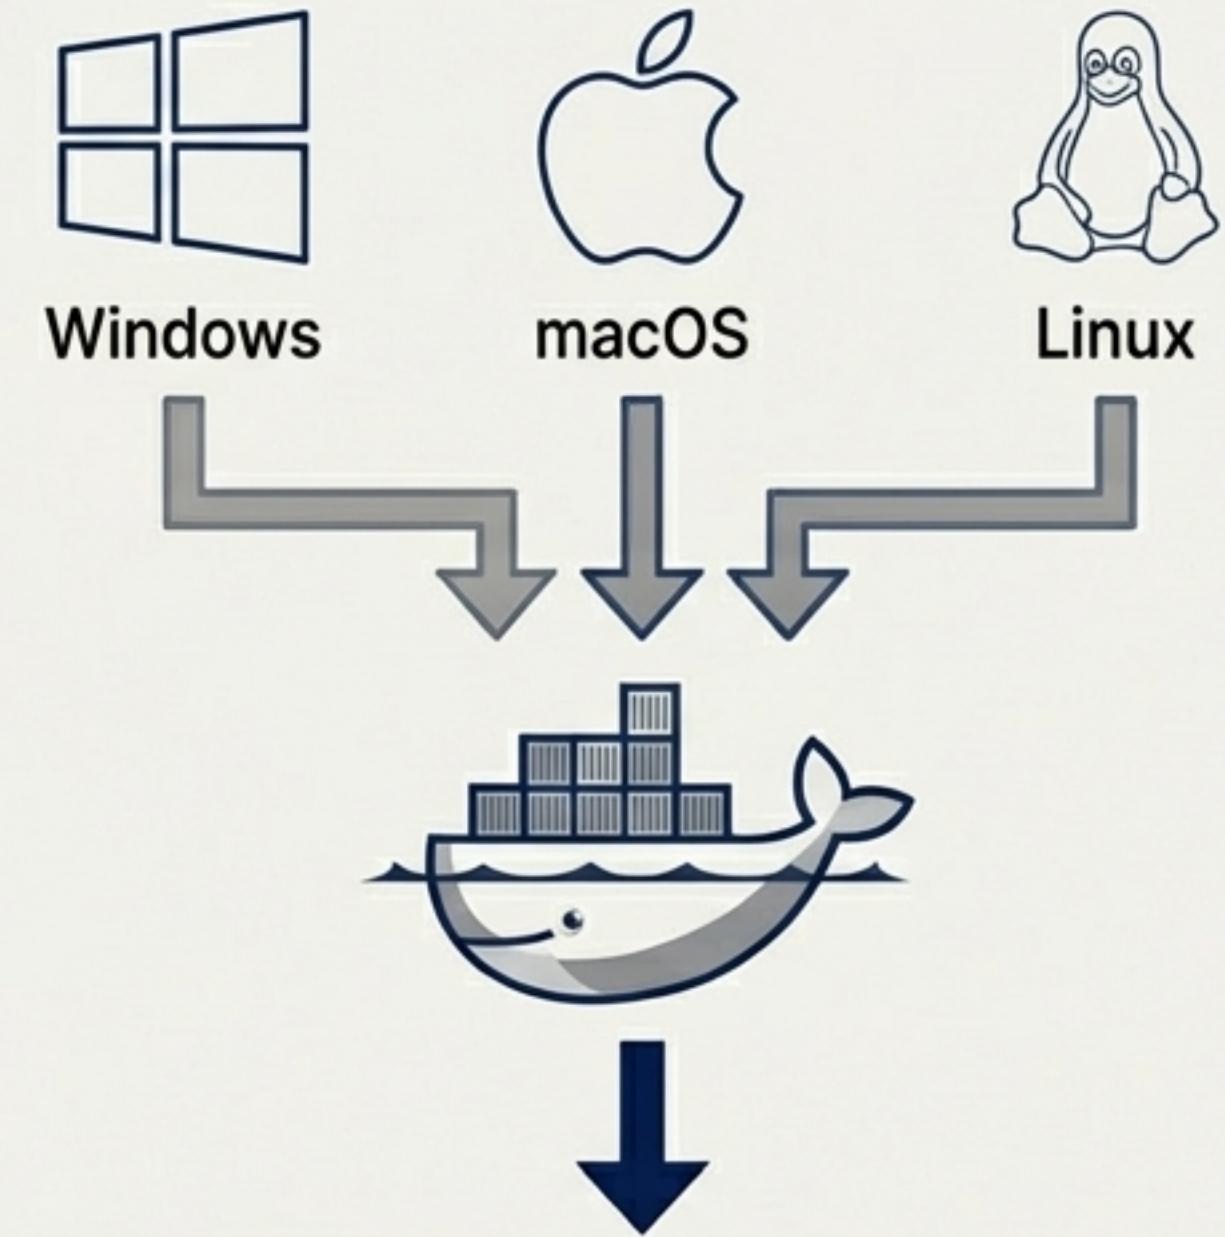
The Problem It Solves

ROS is tightly bound to a particular operating system, usually a specific version of Ubuntu. Setting it up manually can mean fighting with dependencies and versions.

The Docker Solution

- **Consistency:** Guarantees an identical, pre-configured environment on any host OS (Windows, macOS, or Linux).
- **Simplicity:** Avoids manual installation steps and dependency conflicts.
- **Isolation:** Keeps your host system clean. All ROS dependencies live inside the container.

Key Takeaway: We will use a Docker image that has everything pre-installed. This ensures you can follow along with a consistent experience, regardless of your machine.



**Consistent ROS 2 +
Ubuntu Environment**

Launch Sequence: Your First ROS Environment in 3 Steps



1 Install Docker Desktop

Download and install the official Docker Desktop application for your operating system. If on Windows, ensure the WSL 2 backend is enabled.



2 Download the Course Repository

Get the project files from GitHub, which include the Dockerfile needed to build the environment.



3 Build & Run the Container

Navigate to the repository folder in your terminal and run two commands.

```
# 1. Build the image  
docker build -t intro_ros .
```

```
# 2. Run the container  
docker run --rm -it -p 3000:3000 ... intro_ros
```

Note: The run command is simplified for clarity.

Your Journey Begins

- ✓ ROS is the standard **middleware** for robotics, solving common problems.
- ✓ Its architecture is built on **Nodes** communicating via **Topics and Services**.
- ✓ You can use **Python** for rapid development and **C++** for performance.
- ✓ You are now ready to explore with a pre-configured **Docker environment**.

The Destination

The concepts covered in this series will enable you to build real-world applications. Our final demo connects ROS running on a Raspberry Pi to a microcontroller, creating a robot that uses computer vision to find and follow a ball.

